

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.neural_network import MLPRegressor
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
```

```
In [2]: df2 = pd.read_csv('task2.csv')
df2.head()
```

Out[2]:


	x1	x2	y
0	0.996321	2.901429	-61.310885
1	3.855952	2.197317	64.918210
2	-0.751851	1.311989	64.317646
3	-1.535331	2.732352	53.659943
4	2.808920	2.416145	120.558418

The dataset consists of continuous numbers without distinct classes/categories. Hence, the problem represents to be of Regression .

```
In [3]: X = df2.iloc[:, 0:2]
y = df2.iloc[:, 2]
```

```
In [4]: pd.plotting.scatter_matrix(df2)
```

```
Out[4]: array([[<Axes: xlabel='x1', ylabel='x1'>,  
  <Axes: xlabel='x2', ylabel='x1'>,  
  <Axes: xlabel='y', ylabel='x1'>],  
  [<Axes: xlabel='x1', ylabel='x2'>,  
  <Axes: xlabel='x2', ylabel='x2'>,  
  <Axes: xlabel='y', ylabel='x2'>],  
  [<Axes: xlabel='x1', ylabel='y'>, <Axes: xlabel='x2', ylabel='y'>,  
  <Axes: xlabel='y', ylabel='y'>]], dtype=object)
```



It is observed that there exists a correlation between the dependent variable (y) and each of the independent variables (x1 & x2), while the independent variables themselves do not appear to be correlated with each other.

```
In [5]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state = 42)
```

```
In [6]: # After working with multiple values of learning rates, this is the best list of values found.
```

```
learning_rates = [0.3,0.31,0.32,0.33,0.34,0.35,0.36,0.37,0.38,0.39,0.40,0.41,0.42,0.43,0.44,0.45,0.46,0.47,0.48,0.49,
```

```
In [7]: training_scores = []
testing_scores = []
mse_scores = []

for learning_rate in learning_rates:

    # Selecting a Non-Linear model (Neural Network - MLPRegressor) for making predictions.
    reg2 = MLPRegressor(learning_rate='constant', learning_rate_init=learning_rate,
                        hidden_layer_sizes = (100), max_iter = 1000, solver = 'adam', random_state=42)
    reg2.fit(X_train, y_train)

    r2_train_score = reg2.score(X_train, y_train)
    r2_test_score = reg2.score(X_test, y_test)

    y_pred = reg2.predict(X_test)

    mse = mean_squared_error(y_test, y_pred)

    mse_scores.append(mse)
    training_scores.append(r2_train_score)
    testing_scores.append(r2_test_score)

# Checking for model convergence with an if-else statement; if not converged, we increase 'max_iter'.
if reg2.n_iter_ < reg2.max_iter:
    print("The model has converged.")
else:
    print("The model has Not converged.")
```

The model has converged.

```
In [8]: print('Coefficient of Determination of Train Set:\n')
print(training_scores)
print()
print('Coefficient of Determination of Test Set:\n')
print(testing_scores)
print()
print('MSE:\n')
print(mse_scores)
```

Coefficient of Determination of Train Set:

[0.7730642052912683, 0.7677638472768495, 0.7724642933581752, 0.7619249160886976, 0.8567104757916885, 0.7894969488796887, 0.8095611119337671, 0.7592192818963264, 0.7757107478716891, 0.7539608646496643, 0.7962787455495366, 0.7615048381993054, 0.733351543175701, 0.7490123854269676, 0.7057243373677201, 0.757411015532255, 0.799655056678284, 0.7252246233227142, 0.6843151324488479, 0.7141793350228858, 0.7128471241402085]


Coefficient of Determination of Test Set:

[0.7410613438333771, 0.743816620637574, 0.7481623750190556, 0.732828051830944, 0.8136403449058549, 0.7562471067533741, 0.7964986609931246, 0.7345622010280182, 0.7410543169927741, 0.7362630226057589, 0.7427911569765824, 0.7416370032719088, 0.6927799469013534, 0.7116160163508597, 0.6333530872398101, 0.7437452128969304, 0.7242734796369141, 0.6942357483923105, 0.6659018910871475, 0.650859546718688, 0.6538554687518061]

MSE:


[1164.4135790452365, 1152.0234563330414, 1132.481162077307, 1201.4376261682298, 838.035217260776, 1096.1251712270237, 915.1191483013409, 1193.6393819699429, 1164.4451778367688, 1185.9910077566306, 1156.6348335192572, 1161.824913154035, 1381.5287639125938, 1296.825400700842, 1648.7636502533533, 1152.3445676884696, 1239.9064286758069, 1374.9822132756196, 1502.3958975871337, 1570.039370765731, 1556.5670976460356]

```
In [9]: plt.plot(learning_rates, training_scores, label='Training Score', marker = 'o', linewidth = 1, color = 'red')
plt.plot(learning_rates, testing_scores, label='Testing Score', marker = 'o', linewidth = 1, color = 'green')
plt.xlabel('Learning Rate')
plt.ylabel('R2 Score')
plt.title('R2 Score vs Learning Rate')
plt.grid()
plt.legend()
plt.show()
```



The graph above shows us which Learning Rate gives us the highest R2-Score for both the Training Data and Testing Data .

```
In [10]: plt.plot(learning_rates, mse_scores, label='MSE Score', marker = 'o', linewidth = 1, color = 'red')
plt.xlabel('Learning Rate')
plt.ylabel('MSE Score')
plt.title('MSE Score vs Learning Rate')
plt.grid()
plt.legend()
plt.show()
```



The graph above shows us which learning rate gives us the Lowest MSE Score .

```
In [11]: # Selecting the best learning rate based on maximum test R2-Score achieved.
blr = learning_rates[np.argmax(testing_scores)]
print("Best Learning Rate:", blr)
```

Best Learning Rate: 0.34

```
In [12]: # Selecting the best learning rate based on minimum MSE score achieved.
blr = learning_rates[np.argmin(mse_scores)]
print(f"Best Learning Rate:", blr)
```

Best Learning Rate: 0.34

```
In [13]: # Constructing a Best Regression Model using the identified Best Learning Rate.
```

```
best_reg2 = MLPRegressor(learning_rate='constant', learning_rate_init=blr, hidden_layer_sizes=(100), solver = 'adam',
best_reg2.fit(X_train, y_train)
```

Out[13]:

```
MLPRegressor
MLPRegressor(hidden_layer_sizes=100, learning_rate_init=0.34, random_state=42)
```

```
In [14]: mse_train = mean_squared_error(y_train, best_reg2.predict(X_train))
mse_test = mean_squared_error(y_test, best_reg2.predict(X_test))

print("Mean Squared Error on Training Set:", mse_train)
print("Mean Squared Error on Test Set:", mse_test)
```

Mean Squared Error on Training Set: 810.8100968899438  
Mean Squared Error on Test Set: 838.035217260776

After testing different learning rates, the rate of 0.34 gave the lowest MSE Scores : 810.810 for Training and 838.035 for Testing .

```
In [15]: print('Generalization Error:', mse_test - mse_train)
```

Generalization Error: 27.2251203708322

A Generalization Error of 27.22 indicates that our model predicts with High Accuracy due to its low positive value.